



# Welcome to TensorFlow 2.0

Josh Gordon ([twitter.com/random\\_forests](https://twitter.com/random_forests))

Machine Learning Tokyo



# TensorFlow

## An open source Deep Learning library

- Released by Google in 2015
- **>1800** contributors worldwide

## TensorFlow 2.0

- **Easier to use**
- Code styles for beginners and experts
- Alpha released in March, 2019



# Topics

## For beginners and experts

- Keras Sequential
- Keras Subclassing
- Built-in vs custom training loops

## Beyond “Hello World”

- Tutorials for Deep Dream, GANs, Machine Translation

## Under the hood

- AutoGraph and tf.function
- TF2 vs TF1

## Learning more

- Book recommendations



# What exactly is TensorFlow?

- And, what problems are Deep Learning libraries trying to solve?



# Why is Python popular for scientific computing?

—



# Ballpark benchmarks

About how much slower is Python than C?



# Ballpark benchmarks

**About how much slower is Python than C?**

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)



# Ballpark benchmarks

## About how much slower is Python than C?

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)

## Python is a great choice for scientific computing

- Why?





# Ballpark benchmarks

## About how much slower is Python than C?

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)

## Python is a great choice for scientific computing

- Why?

## NumPy

- C performance, Python ease of use



# TensorFlow is basically

## NumPy

- + GPU / TPU support
- + AutoDiff
- + Utilities to help you write neural networks (layers, optimizers)

## TensorFlow

- A C++ engine to accelerate code written in Python.
- **Bonus:** your program is compiled to a graph that can run on devices **without a Python interpreter** (phones, web browsers)

# You can use TF 2.0 like NumPy

```
import tensorflow as tf # Assuming TF 2.0 is installed

a = tf.constant([[1, 2],[3, 4]])
b = tf.matmul(a, a)

print(b)
# tf.Tensor( [[ 7 10] [15 22]], shape=(2, 2), dtype=int32)

print(type(b.numpy()))
# <class 'numpy.ndarray'>
```



# Exercise 1

## Goals

- Install TensorFlow 2.0
- Introduce Colab
- Introduce gradient descent

## Visit

- [bit.ly/tf-ws1](https://bit.ly/tf-ws1)



# For beginners and experts

# For beginners

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# TF 1.x

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# TF 2.0

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```







# Keras and tf.keras

In my view, the **clearest Deep Learning library** that exists today.

- For fast prototyping, advanced research, and production.

**keras.io** = reference implementation

- `import keras`

**tf.keras** = TensorFlow's implementation (a superset, built-in to TF, no need to install Keras separately)

- `from tensorflow import keras`

# For experts

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__(name='my_model')  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes, activation='sigmoid')  
  
    def call(self, inputs):  
        # Define your forward pass here,  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```



# What's the difference?



# Symbolic vs Imperative APIs

## Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**



# Symbolic vs Imperative APIs

## Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**

## Imperative (Keras Subclassing)

- Your model is Python bytecode
- Complete flexibility and control
- Harder to debug / **harder to maintain**

# Use a built-in training loop...

```
model.fit(x_train, y_train, epochs=5)
```

# Or define your own

```
model = MyModel()
```

```
with tf.GradientTape() as tape:  
    logits = model(images)  
    loss_value = loss(logits, labels)
```

```
grads = tape.gradient(loss_value, model.trainable_variables)  
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```



# TensorBoard

```
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
```

```
model.fit(  
    x_train, y_train, epochs=5,  
    validation_data=[x_test, y_test],  
    callbacks=[tb_callback])
```

- ☐ Show data download links
- ☒ Ignore outliers in chart scaling

Tooltip sorting method: **default**

Smoothing



Horizontal Axis

STEP

RELATIVE

WALL

Runs

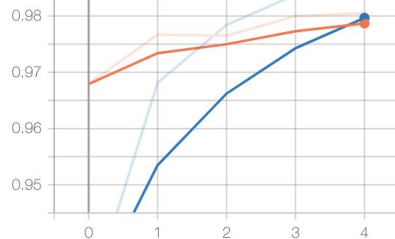
Write a regex to filter runs

- ☒ 20190227-033014/test
- ☒ 20190227-033014/train

Filter tags (regular expressions supported)

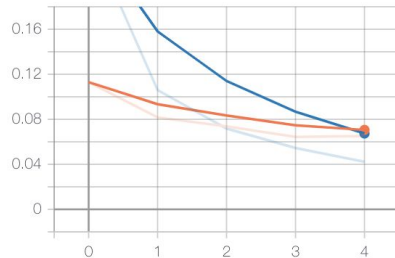
accuracy

accuracy  
tag: accuracy

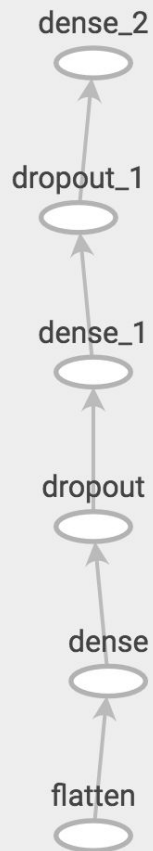


loss

loss  
tag: loss



sequential





# Beyond Hello World



# A few of my favorites

- Machine Translation
- Image Captioning (incidentally, the decoder is similar!)
- DCGan and Pix2Pix



# The docs are code

## Tutorials on [tf.org/alpha](https://tf.org/alpha) are

- Backed by a Jupyter Notebook
- Can be run directly in Colab

## They automatically


- Install the right TensorFlow version
- Download a dataset
- Train a model
- Show you the result

[tensorflow.org/alpha/tutorials/text/image\\_captioning](https://tensorflow.org/alpha/tutorials/text/image_captioning)

TensorFlow > Learn > TensorFlow Core > TF 2.0 Alpha

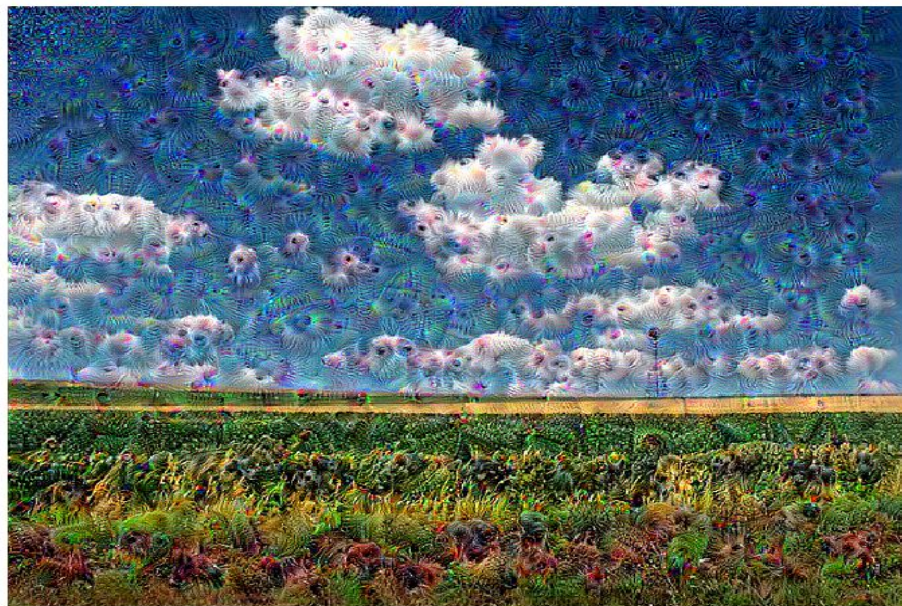
## Image Captioning with Attention

 Run in Google Colab

 View source on GitHub

Given an image like the below, our goal is to generate a caption, such as "a surfer riding on a wave".

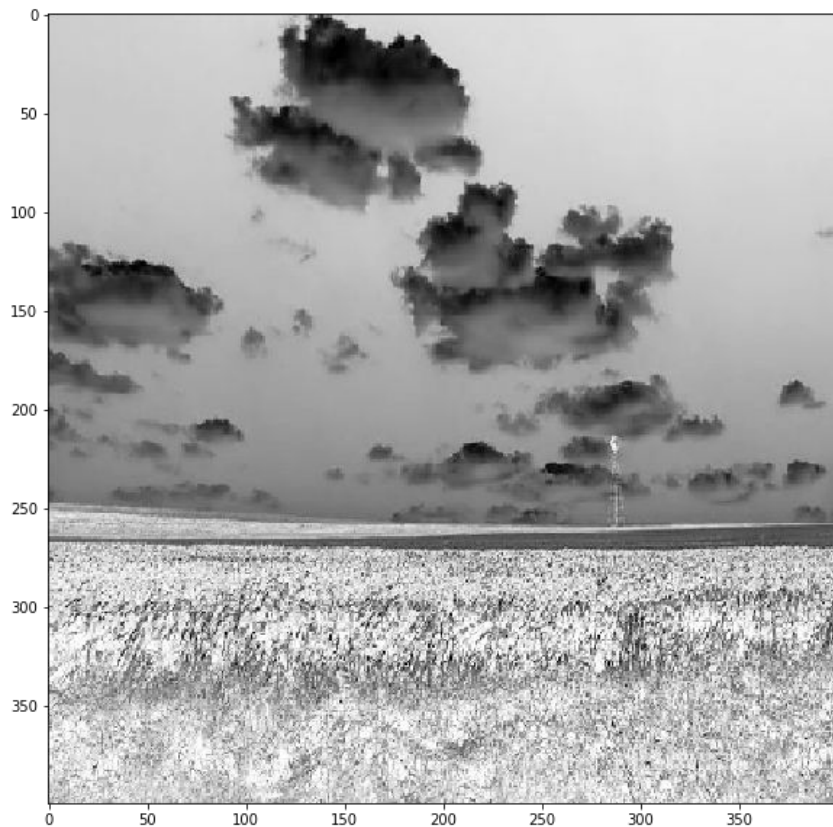




<https://github.com/random-forests/applied-dl/blob/master/examples/9-deep-dream-minimal.ipynb>

# Code walkthrough





<https://github.com/random-forests/applied-dl/blob/master/examples/9-image-colorization.ipynb>





# Is anyone bilingual? Trilingual?

**When translating, do you...**

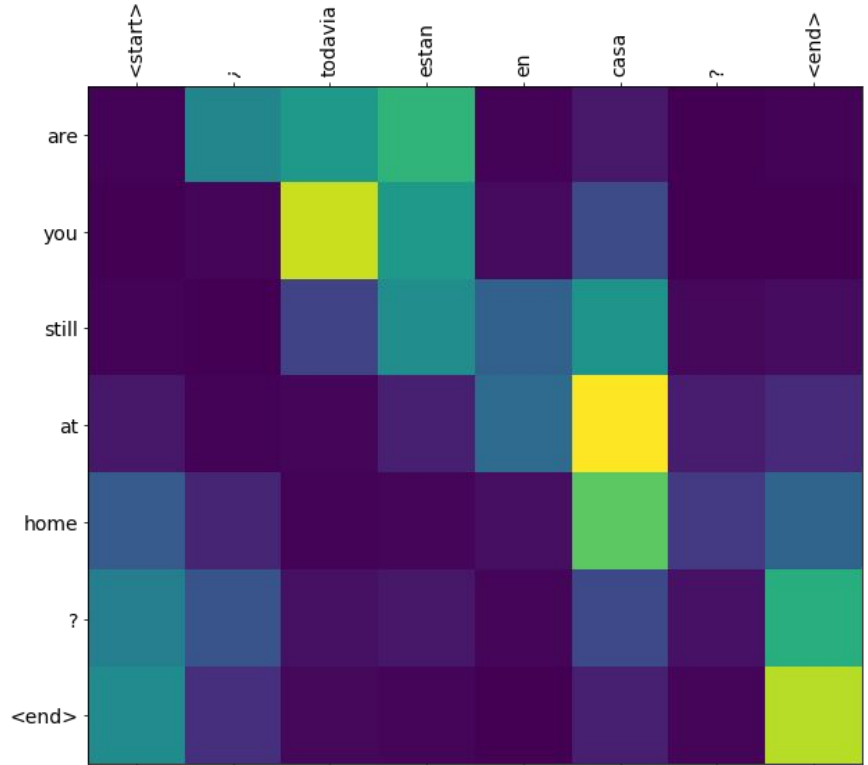
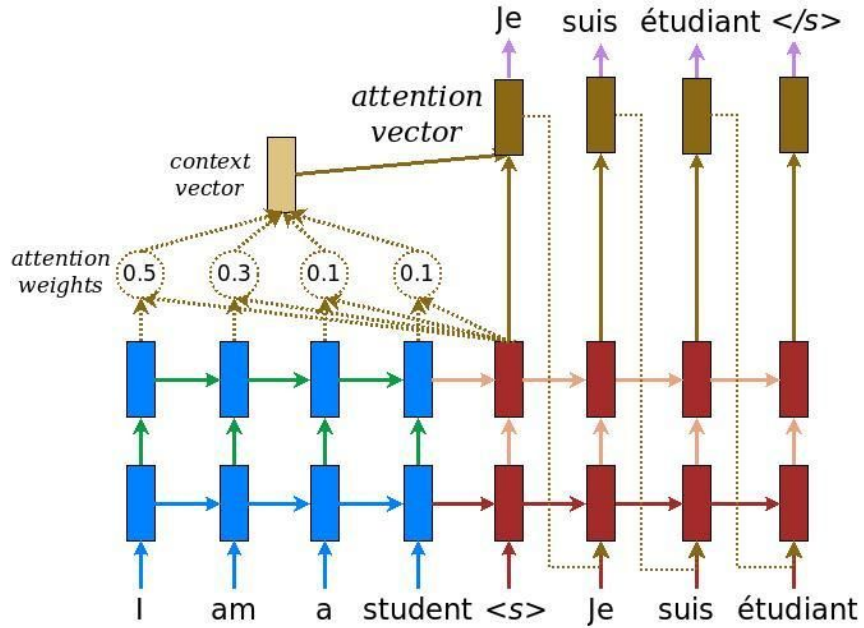
- Go directly from source -> target
- Or, go from source -> **intermediate representation** -> target.



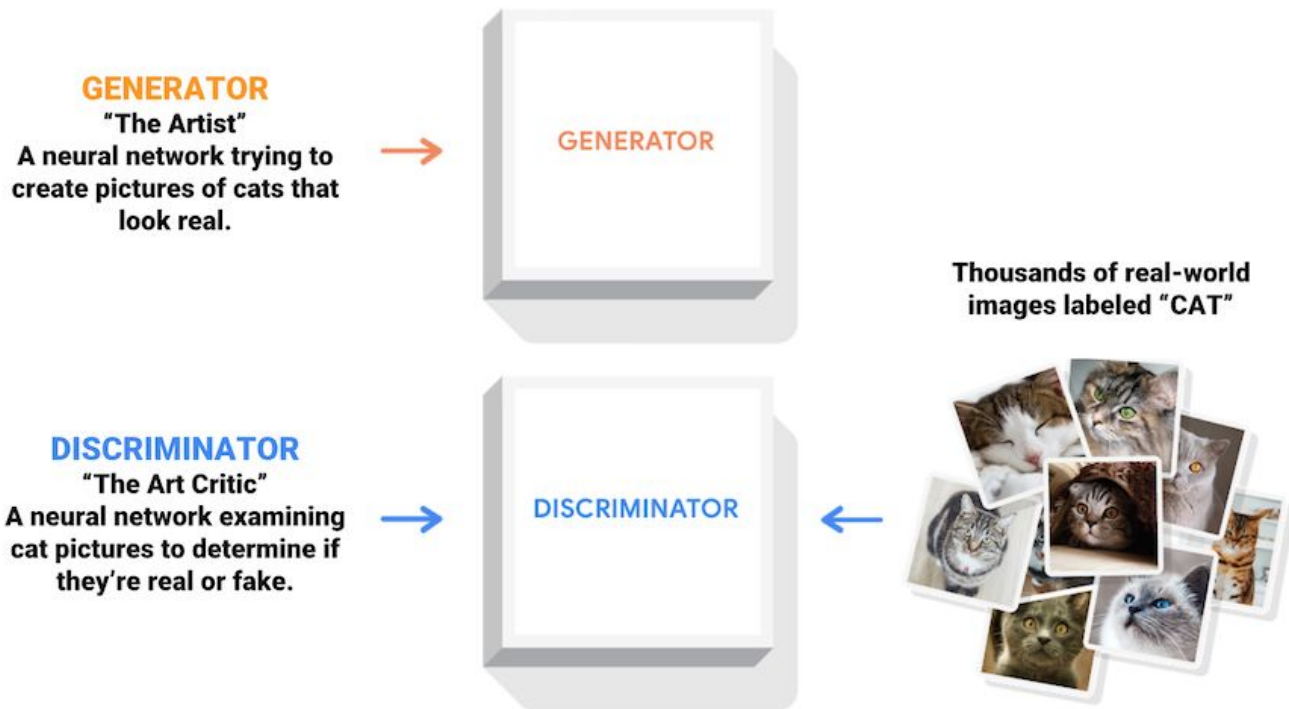
# Machine translation tutorials

- [Hello world](#) (seq2seq), trains in about a minute.
- [Neural Machine Translation with Attention](#)
- [Transformer](#)

P.S., isn't 2019 cool? It's **amazing** this is possible.

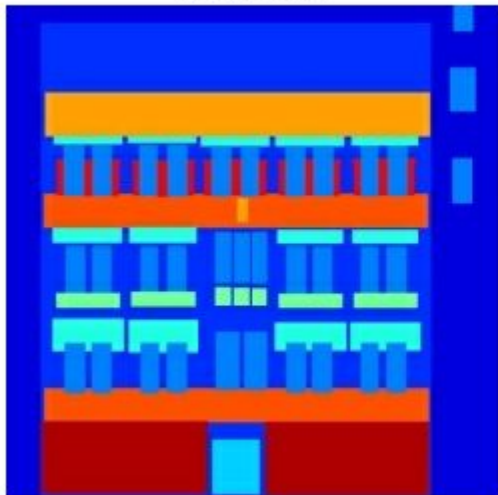


[https://www.tensorflow.org/alpha/tutorials/sequences/nmt\\_with\\_attention](https://www.tensorflow.org/alpha/tutorials/sequences/nmt_with_attention)



<https://www.tensorflow.org/alpha/tutorials/generative/dcgan>

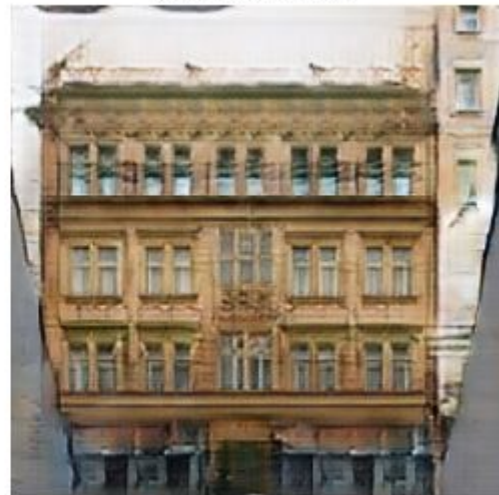
Input Image



Ground Truth



Predicted Image



<https://www.tensorflow.org/alpha/tutorials/generative/pix2pix>



Prediction Caption: the person is riding a surfboard in the ocean <end>

[https://www.tensorflow.org/alpha/tutorials/sequences/image\\_captioning](https://www.tensorflow.org/alpha/tutorials/sequences/image_captioning)



# Under the hood

# Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
def fn(input, state):  
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2  
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```



# Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
@tf.function
```

```
def fn(input, state):
```

```
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
```

```
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

```
timeit.timeit(lambda: fn(input, state), number=10) # 0.004
```

# AutoGraph makes this possible

```
@tf.function
```

```
def f(x):
```

```
    while tf.reduce_sum(x) > 1:
```

```
        x = tf.tanh(x)
```

```
    return x
```

```
# you never need to run this (unless curious)
```

```
print(tf.autograph.to_code(f))
```

# Generated code

```
def tf__f(x):  
    def loop_test(x_1):  
        with ag__.function_scope('loop_test'):  
            return ag__.gt(tf.reduce_sum(x_1), 1)  
    def loop_body(x_1):  
        with ag__.function_scope('loop_body'):  
            with ag__.utils.control_dependency_on_returns(tf.print(x_1)):  
                tf_1, x = ag__.utils.alias_tensors(tf, x_1)  
                x = tf_1.tanh(x)  
            return x,  
    x = ag__.while_stmt(loop_test, loop_body, (x,), (tf,))  
    return x
```

# Going big: tf.distribute.Strategy

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Dense(64, input_shape=[10]),  
    tf.keras.layers.Dense(64, activation='relu'),  
    tf.keras.layers.Dense(10, activation='softmax')])  
  
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

# Going big: Multi-GPU

```
strategy = tf.distribute.MirroredStrategy()

with strategy.scope():
    model = tf.keras.models.Sequential([
        tf.keras.layers.Dense(64, input_shape=[10]),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax')])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```



# What's different between TF1 and TF2?

## Removed

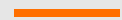
- `session.run`
- `tf.control_dependencies`
- `tf.global_variables_initializer`
- `tf.cond`, `tf.while_loop`

## Added

- `tf.function`, `AutoGraph`



# TensorFlow.js



# Demo #1

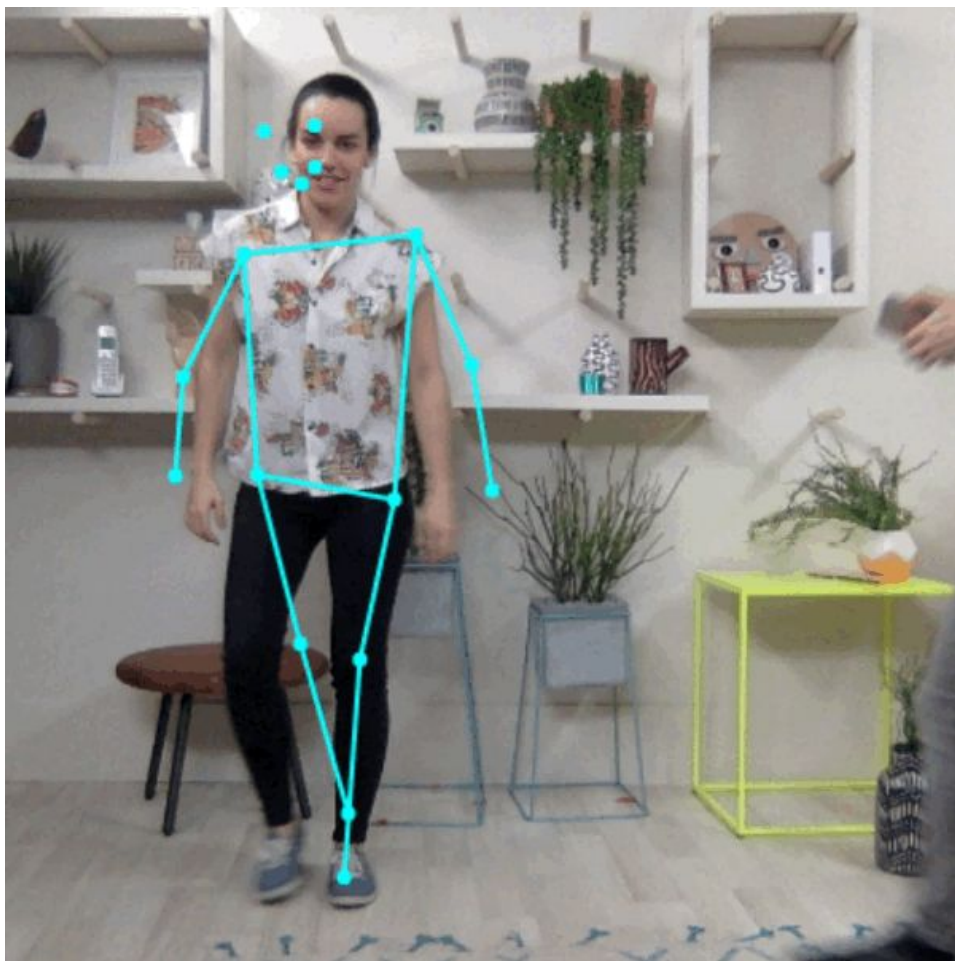
PoseNet





# PoseNet

[bit.ly/pose-net](https://bit.ly/pose-net)



---

# Demo #2

BodyPix



# BodyPix

[bit.ly/body-pix](https://bit.ly/body-pix)





# Learning more



# Learn more

## Tutorials and guides

- [tensorflow.org/alpha](https://www.tensorflow.org/alpha)

## Books

- [Deep Learning with Python](#)
- Hands-On Machine Learning with Scikit-Learn and TensorFlow (version 2.0 is almost ready)

## Courses

- [Intro to Deep Learning](#) (MIT)
- [Convolutional Neural Networks for Visual Recognition](#) (Stanford)



# tf.thanks!

Josh Gordon ([twitter.com/random\\_forests](https://twitter.com/random_forests))